

Assignment1 Report (Rosie Wang 1806394)

- Which regularization worked better (L1 vs L2)?

L2 works better in my code(L1_lambda=0.001) with less training loss in the same epoch.

```
⇒ Training with L2 Regularization:  
Epoch [1/10], Loss with L2: 0.2581  
Epoch [2/10], Loss with L2: 0.2563  
Epoch [3/10], Loss with L2: 0.2554  
Epoch [4/10], Loss with L2: 0.2545  
Epoch [5/10], Loss with L2: 0.2517  
Epoch [6/10], Loss with L2: 0.2509  
Epoch [7/10], Loss with L2: 0.2508  
Epoch [8/10], Loss with L2: 0.2483  
Epoch [9/10], Loss with L2: 0.2494  
Epoch [10/10], Loss with L2: 0.2473
```

```
Training with Manual L1 Regularization:  
Epoch [1/10], Loss with L1: 0.6286  
Epoch [2/10], Loss with L1: 0.5148  
Epoch [3/10], Loss with L1: 0.4735  
Epoch [4/10], Loss with L1: 0.4573  
Epoch [5/10], Loss with L1: 0.4510  
Epoch [6/10], Loss with L1: 0.4474  
Epoch [7/10], Loss with L1: 0.4465  
Epoch [8/10], Loss with L1: 0.4448  
Epoch [9/10], Loss with L1: 0.4437  
Epoch [10/10], Loss with L1: 0.4450
```

- Which optimizer was more effective (SGD vs Adam)?

```
print('Validation Accuracy: %.100 %'
```


```
⇒ Epoch [1/10], Loss: 0.2489  
Validation Accuracy: 91.43%  
Epoch [2/10], Loss: 0.2485  
Validation Accuracy: 91.13%  
Epoch [3/10], Loss: 0.2473  
Validation Accuracy: 91.75%  
Epoch [4/10], Loss: 0.2471  
Validation Accuracy: 90.50%  
Epoch [5/10], Loss: 0.2453  
Validation Accuracy: 91.14%  
Epoch [6/10], Loss: 0.2436  
Validation Accuracy: 91.52%  
Epoch [7/10], Loss: 0.2434  
Validation Accuracy: 91.68%  
Epoch [8/10], Loss: 0.2431  
Validation Accuracy: 91.29%  
Epoch [9/10], Loss: 0.2415  
Validation Accuracy: 91.06%  
Epoch [10/10], Loss: 0.2421  
Validation Accuracy: 91.08%
```

```
⇒ Training with Adam Optimizer:  
Epoch [1/10], Loss with Adam: 0.2383  
Validation Accuracy: 91.39%  
Epoch [2/10], Loss with Adam: 0.2370  
Validation Accuracy: 91.30%  
Epoch [3/10], Loss with Adam: 0.2352  
Validation Accuracy: 91.55%  
Epoch [4/10], Loss with Adam: 0.2347  
Validation Accuracy: 91.32%  
Epoch [5/10], Loss with Adam: 0.2346  
Validation Accuracy: 91.56%  
Epoch [6/10], Loss with Adam: 0.2344  
Validation Accuracy: 91.48%  
Epoch [7/10], Loss with Adam: 0.2337  
Validation Accuracy: 91.64%  
Epoch [8/10], Loss with Adam: 0.2325  
Validation Accuracy: 90.68%  
Epoch [9/10], Loss with Adam: 0.2319  
Validation Accuracy: 91.41%  
Epoch [10/10], Loss with Adam: 0.2313  
Validation Accuracy: 91.40%
```

According to the output, we can tell Adam (figure 2) has the higher validation accuracy compared to the SGD's validation accuracy in the same epoch. So Adam is more effective.

• What does the confusion matrix reveal?


confusion matrix shows how is the performance the model hold. In my case, the confusion matrix reveals that logistic regression achieves high accuracy on simple digits like **0, 1, 7, and 9**, but struggles with visually similar pairs such as **2 vs 8, 3 vs 5/8, and 4 vs 9**. This highlights the model's limitation in capturing fine-grained local features compared to more advanced models like CNNs.



- **How does performance degrade with reduced or noisy data?**

Performance(from low to high): noisy data training < 10% training < 25% training < 50% training < 100% training.

Training with 10.0% of the training data...
Validation Accuracy with 6000 training samples: 90.21%
Training with 25.0% of the training data...
Validation Accuracy with 15000 training samples: 90.86%
Training with 50.0% of the training data...
Validation Accuracy with 30000 training samples: 92.39%



- **How did Gemini help (include sample prompts + modifications)?**

Gemini helped me understand the code and some specific and common syntax in pytorch and CNN. Also, it helps debug with the process of CNN and modifying with the hyperparameter etc.